Aleem Sattar

Software Engineer — Data Platform

Dallas–Fort Worth, TX | aleem.sattar@gmail.com | linkedin.com/in/aleem-s-381985147 | github.com/Aleem89

SUMMARY

Software engineer who builds and operates production data pipelines on Python, PostgreSQL, and Kubernetes for a multi-brand franchise analytics warehouse: REST API ingestion, dimensional modeling, Superset/Streamlit delivery, and metric reconciliation with stakeholders. Apache Superset contributor; prior SaaS support career brings a debugging-first, customer-facing instinct.

SKILLS

Languages: Python, SQL, Bash, JavaScript, TypeScript, HTML, CSS, PHP

Data & Analytics: PostgreSQL (views, MVs, dimensional modeling), SQLAlchemy, Alembic, pandas, Airbyte, Apache Superset, OpenMetadata, Cube.js, Streamlit

Platform & DevOps: Docker, Kubernetes (AKS), Helm, GitHub Actions, Sealed Secrets, Azure CLI, Railway, Linux

Backend & Integration: FastAPI, REST API clients, webhooks, async I/O, Pydantic, pytest, Click CLIs

Identity & Access: OIDC / OAuth2, Microsoft Entra ID, Keycloak, OAuth2-proxy, RBAC, row-level security, PingOne

Practices: Agile, PR code review, runbooks, data-quality reconciliation, AI-assisted development (Claude Code, MCP)

EXPERIENCE

Savantly LLC — Software Engineer | Remote (DFW) **Jul 2025 – Present**

Consultancy; embedded on a multi-brand franchisor's data platform and web properties with a senior engineer and contractors.

- Built and operate 8 production data pipelines (Python sync services + Airbyte syncs) into a PostgreSQL warehouse with idempotent snapshot loads, replacing ~30 hrs/month of manual spreadsheet consolidation.
- Designed the dimensional model and 60+ SQL views/MVs behind franchise KPIs (same-store sales, TTM revenue, territory counts) for 200+ locations, consumed by Superset and Streamlit; 25+ Alembic migrations.
- Deployed every pipeline and the analytics app to AKS via Helm CronJobs, GitHub Actions CI/CD, and sealed secrets; moved the scorecard app from Railway onto AKS for one deploy path and ~10-minute releases.
- Cut scorecard page loads from 7.3 s to 0.2 s (Overview) and 1.9 s to 35 ms (master view), a ~97% reduction, by precomputing revenue and GA4 attribution into materialized views and collapsing duplicate query passes.
- Replaced a truncate-and-reload territory-metrics table with a live SQL view over HubSpot data (−1,070 lines of ETL); client territory counts went from up to a month stale to current, at <10 ms added query cost.
- Led a phased location-mapping cutover — shadow views validated against 4 downstream consumers, deterministic refresh, lock timeouts, written runbook — with zero reporting regressions.
- Own the production Apache Superset platform (31 dashboards, 24 SSO users), avoiding an estimated $10K+/yr in per-seat BI licensing; upgraded 5.0→6.1 with zero downtime via staged PRs and cache warm-up.
- Reconciled KPI definitions (answer rate, same-store sales, territory counts) with client finance and marketing stakeholders, fixing reported discrepancies within a business day; published an OpenMetadata glossary.

Omnicell — Technical Support Engineer (SaaS) | Fort Worth, TX **Apr 2024 – Apr 2025**

- Technical contact for 130+ pharmacy SaaS accounts; closed 95% of Salesforce cases within the 24-hour SLA.
- Streamlined SFTP data-exchange onboarding, cutting per-account setup time 20% (about a day each).
- Reproduced and escalated defects to engineering in Jira; analysis lifted compliant stores 40% in 6 months.

Santander Consumer USA — Team Lead, Platinum Team | North Richland Hills, TX **Jun 2019 – Apr 2023**

- Escalation lead for regulated accounts (bankruptcy, SCRA, credit disputes); tracked issues in ServiceNow.
- Sustained 95%+ CSAT and 90%+ first-contact resolution; mentored junior representatives.

EDUCATION & OPEN SOURCE

University of Texas at Dallas — B.A., Biology, cum laude Richardson, TX

Apache Superset — Contributor: caught a Pivot Table regression in the 7.0.0 RC, filed issue #45405, flagged it on the release vote, and shipped the fix (PR #45408, merged in <24 h); earlier docs PR #44108.